

# Project 2

Rachael Stryker, Tom Muhlbauer, and Daniel Lassiter

12/4/2020

## Contents


|          |  |           |
|----------|--|-----------|
| <b>1</b> | <b>Exploratory Data Analysis</b>   | <b>1</b>  |
| 1.1      | Inspecting for trends and seasonality . . . . .                                | 1         |
| 1.2      | Hypotheses from exploratory data analysis: . . . . .                           | 6         |
| <b>2</b> | <b>Building Univariate Time Series Models</b>                                  | <b>6</b>  |
| 2.1      | Modeling trend and season . . . . .  | 6         |
| 2.2      | Modeling the residuals of best trend/season models . . . . .                   | 13        |
| <b>3</b> | <b>Forecasting</b>   | <b>23</b> |
| 3.1      | Combining Season/trend and residuals forecasts and computing MSE . . . . .     | 27        |
| 3.2      | Plotting forecasts and MSE . . . . .   | 27        |
| <b>4</b> | <b>Simulation</b>  | <b>28</b> |
| 4.1      | Simulating one year proceeding our observations . . . . .                      | 28        |
| 4.2      | Comparing trend/season models built from observations vs. simulation . . . . . | 30        |
| 4.3      | Verifying simulation reproduces seasonality . . . . .                          | 31        |
| 4.4      | Verifying simulation reproduces autocorrelation structure . . . . .            | 34        |

## 1 Exploratory Data Analysis

### 1.1 Inspecting for trends and seasonality

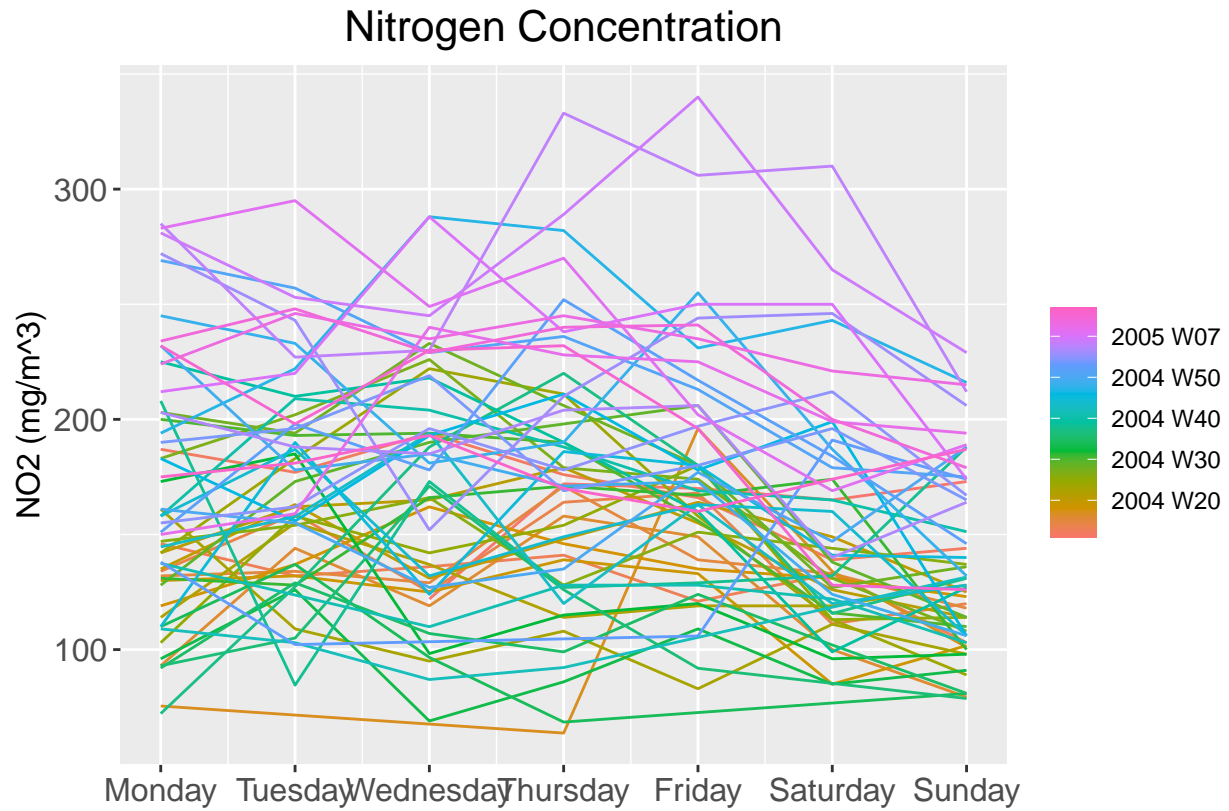
```
# For quickly formatting plots
my_theme <- theme(plot.title = element_text(hjust = 0.5, size = 16),
  axis.title = element_text(size = 12),
  axis.text = element_text(size = 12),
  plot.caption = element_text(size = 12, hjust = 0))
```

```
dailyAQ %>% autoplot(NO2) + ylab("NO2 (mg/m^3)") + xlab("Date") +  
  ggtitle("Nitrogen Concentration") + my_theme
```



The time series plot indicates potential seasonal components and a potential trend component.

```
dailyAQ %>% gg_season(NO2, period = "week") + ylab("NO2 (mg/m^3)") + xlab("") +  
  ggtitle("Nitrogen Concentration") + my_theme
```

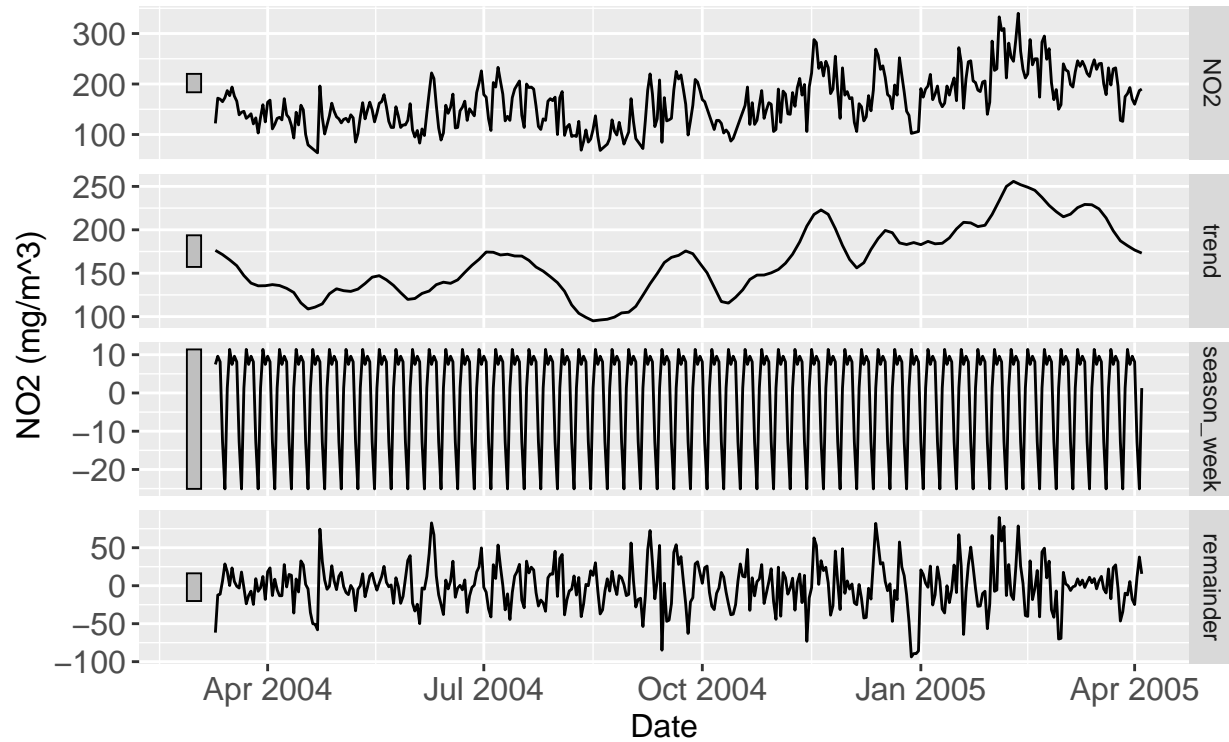


This plot indicates an increasing trend over time.

```
dailyAQ %>% model(STL(NO2 ~ trend(window=21) + season(window = "periodic"),
  robust = TRUE)) %>%
  components() %>%
  autoplot() + my_theme + ylab("NO2 (mg/m^3)")
```

## STL decomposition

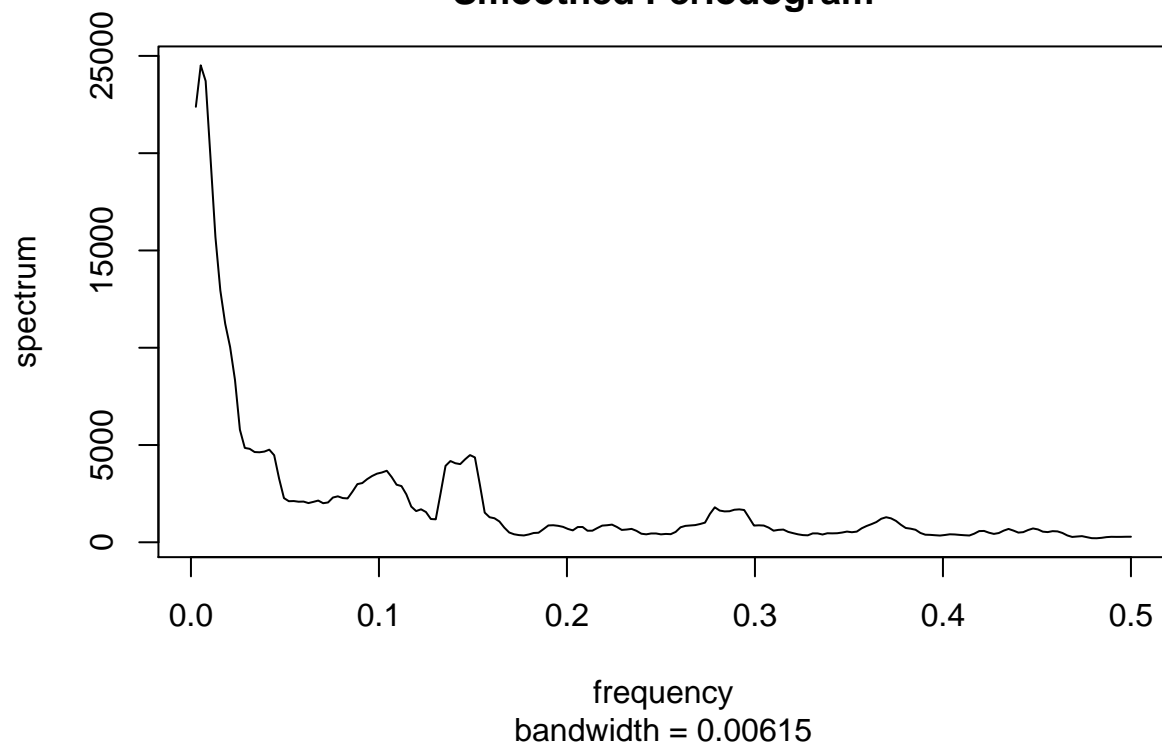
NO2 = trend + season\_week + remainder



This seasonal decomposition chart shows a possibly increasing trend and a weekly seasonal component.

```
NO2.ts <- head(NO2.ts, n = length(NO2.ts) - 8)
pg.NO2 <- spec.pgram(NO2.ts, spans=9, demean=T, log='no')
```

# Series: NO2.ts Smoothed Periodogram



```
spec.NO2 <- data.frame(freq=pg.NO2$freq, spec=pg.NO2$spec)

spec.NO2 <- data.frame(freq=pg.NO2$freq, spec=pg.NO2$spec)%>%
  mutate(period = 1/freq) %>% arrange(desc(spec))

spec.NO2[c(1:17), ]
```

| ##    | freq        | spec      | period     |
|-------|-------------|-----------|------------|
| ## 1  | 0.005208333 | 24514.627 | 192.000000 |
| ## 2  | 0.007812500 | 23711.028 | 128.000000 |
| ## 3  | 0.002604167 | 22385.338 | 384.000000 |
| ## 4  | 0.010416667 | 19749.539 | 96.000000  |
| ## 5  | 0.013020833 | 15692.398 | 76.800000  |
| ## 6  | 0.015625000 | 12945.511 | 64.000000  |
| ## 7  | 0.018229167 | 11232.528 | 54.857143  |
| ## 8  | 0.020833333 | 10035.151 | 48.000000  |
| ## 9  | 0.023437500 | 8340.280  | 42.666667  |
| ## 10 | 0.026041667 | 5785.995  | 38.400000  |
| ## 11 | 0.028645833 | 4849.327  | 34.909091  |
| ## 12 | 0.031250000 | 4803.541  | 32.000000  |
| ## 13 | 0.041666667 | 4760.502  | 24.000000  |
| ## 14 | 0.039062500 | 4662.202  | 25.600000  |
| ## 15 | 0.033854167 | 4638.409  | 29.538462  |
| ## 16 | 0.036458333 | 4627.030  | 27.428571  |
| ## 17 | 0.148437500 | 4482.699  | 6.736842   |

There is a peak at around .15 which corresponds to a seasonality component with a period of 1 week. There is also a possible monthly season and longer seasonal periods are also possible. We decided to model many combinations of these periods.

## 1.2 Hypotheses from exploratory data analysis:

There is an increasing trend in NO2 emissions. This could be due to increasing cars on the road over time. There is also one or more seasonal component.

## 2 Building Univariate Time Series Models

```
# Creating potential predictors
dailyAQ <- dailyAQ %>% mutate(step = 1:length(Date)) %>%
  mutate(day = as.factor(as.character(lubridate::wday(Date, label = TRUE)))) %>%
  mutate(weekend = isWeekend(Date))

# Create train and test sets
dailyAQ_train <- dailyAQ %>% head(n = length(Date) - 8)
dailyAQ_test <- dailyAQ %>% tail(n = 7)
```

### 2.1 Modeling trend and season

Deciding how to model the seasonal component was an iterative process. The models shown below include only those we considered to be reasonable candidates. Note that we decided to use the packages feasts, fable, and fabletools for these analyses so our code looks different than the the in class examples.

```
# Create potential trend/season models
models <- dailyAQ_train %>%
  model(trend = TSLM(NO2 ~ trend()),
        trnd_ssn_w_wkday_dummy = TSLM(NO2 ~ trend() + day),
        trnd_ssn_w_wknd_dummy = TSLM(NO2 ~ trend() + weekend),
        trnd_ssn_wkly = TSLM(NO2 ~ trend() + sin(2*pi*step/7) + cos(2*pi*step/7)),
        trnd_ssn_wkly_mnthly = TSLM(NO2 ~ trend() + sin(2*pi*step/(365.25/12)) +
                                     cos(2*pi*step/(365.25/12)) +
                                     sin(2*pi*step/7) + cos(2*pi*step/7)),
        trnd_ssn_complex_trig = TSLM(NO2 ~ trend() + sin(2*pi*step/192) +
                                     cos(2*pi*step/192) + sin(2*pi*step/128) +
                                     cos(2*pi*step/128) + sin(2*pi*step/76.8) +
                                     cos(2*pi*step/76.8) + sin(2*pi*step/64) +
                                     cos(2*pi*step/64) + sin(2*pi*step/7) +
                                     cos(2*pi*step/7) + sin(2*pi*step/30) +
                                     cos(2*pi*step/30) + sin(2*pi*step/365) +
                                     cos(2*pi*step/365)),
        trnd_ssn_complex_trig_and_dummy = TSLM(NO2 ~ trend() + weekend +
                                                sin(2*pi*step/192) +
                                                cos(2*pi*step/192) + sin(2*pi*step/128) +
                                                cos(2*pi*step/128) + sin(2*pi*step/76.8) +
                                                cos(2*pi*step/76.8) + sin(2*pi*step/64) +
                                                cos(2*pi*step/64) + sin(2*pi*step/7) +
```

```

cos(2*pi*step/7)+sin(2*pi*step/30) +
cos(2*pi*step/30)+sin(2*pi*step/365) +
cos(2*pi*step/365))
)

```

```

# Inspect the trend only model
models %>% dplyr::select(trend) %>% report()
## Series: NO2
## Model: TSLM
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -88.444 -33.593   2.634  28.304 138.425
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 113.76664    4.51323   25.21  <2e-16 ***
## trend()      0.25902    0.02032   12.75  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 44.13 on 382 degrees of freedom
## Multiple R-squared:  0.2985, Adjusted R-squared:  0.2966
## F-statistic: 162.5 on 1 and 382 DF, p-value: < 2.22e-16

```


The coefficient on trend is significant, and plotting the modeled trend line against the original data below seems to verify this. The significance of the trend will change when the seasonal components are modeled but it turns out to be significant in all cases.

```

models %>% dplyr::select(trend) %>% fitted() %>%
  autoplot(color = 'red') + autolayer(dailyAQ_train) +
  ggtitle("Trend Model") + ylab("NO2 mg/m^3") + xlab("") + my_theme
## Plot variable not specified, automatically selected '.vars = .fitted'
## Plot variable not specified, automatically selected '.vars = NO2'

```

## Trend Model



Next, we inspected the reports and plots of the trend season models that used day-of-the-weekend and weekend/not weekend dummy variables.


```
models %>% dplyr::select(trnd_ssn_w_wkday_dummy) %>% report()
## Series: NO2
## Model: TSLM
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -99.28 -26.90   2.88  25.05 128.38
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 123.4759    6.8287   18.082 < 2e-16 ***
## trend()      0.2600     0.0195   13.333 < 2e-16 ***
## dayMon      -7.6160     8.0767   -0.943  0.34630
## daySat     -22.7403     8.0765   -2.816  0.00512 **
## daySun     -36.9244     8.0765   -4.572 6.57e-06 ***
## dayThu       1.9389     8.0765    0.240  0.81041
## dayTue      -1.8427     8.1138   -0.227  0.82046
## dayWed      -1.9503     8.0765   -0.241  0.80931
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 42.35 on 376 degrees of freedom
## Multiple R-squared:  0.3641, Adjusted R-squared:  0.3523
```



```
## F-statistic: 30.75 on 7 and 376 DF, p-value: < 2.22e-16

models %>% dplyr::select(trnd_ssn_w_wkday_dummy) %>% fitted() %>%
  autoplot(color = 'red') + autolayer(dailyAQ_train) +
  ggtitle("Trend-Season Model with Day-of-the-week Dummy Variables") +
  ylab("NO2 mg/m^3") + xlab("") + my_theme
## Plot variable not specified, automatically selected '.vars = .fitted'
## Plot variable not specified, automatically selected '.vars = NO2'
```

## Trend-Season Model with Day-of-the-week Dummy Variable



This model used a dummy variable for each day of the week. Friday is the base case. Only Saturday and Sunday were significantly different from the base case. NO<sub>2</sub> concentrations on these days are lower. Perhaps traffic is heavier during the weekdays because of commuters. It is also notable that the trend term is still significant once the seasonality component has been accounted for.


The insignificance of the weekday variables led us to attempt to model the seasonal component with a single dummy variable indicating whether the date was on a weekday or weekend.

```
models %>% dplyr::select(trnd_ssn_w_wknd_dummy) %>% report()
## Series: NO2
## Model: TSLM
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -96.440  -27.579   2.535   26.753  130.327
##
## Coefficients:
```

```
##               Estimate Std. Error t value Pr(>|t|)
## (Intercept) 121.64424    4.53395  26.830 < 2e-16 ***
## trend()      0.25967     0.01949  13.324 < 2e-16 ***
## weekendTRUE -27.93766     4.77846  -5.847 1.08e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 42.33 on 381 degrees of freedom
## Multiple R-squared:  0.3562, Adjusted R-squared:  0.3529
## F-statistic: 105.4 on 2 and 381 DF, p-value: < 2.22e-16

models %>% dplyr::select(trnd_ssn_w_knd_dummy) %>% fitted() %>%
  autoplot(color = 'red') + autolayer(dailyAQ_train) +
  ggtitle("Trend-Season Model with Dummy Variables") + ylab("NO2 mg/m^3") +
  xlab("") + my_theme
## Plot variable not specified, automatically selected '.vars = .fitted'
## Plot variable not specified, automatically selected '.vars = NO2'
```

## Trend-Season Model with Dummy Variables



All coefficients are significant indicating this is a strong candidate model.

Next we inspected the model reports and plot of the trend season models that relied on trigonometric functions. We show this all in a single graph and table to save space. Note that the model 'trnd\_ssn\_complex\_trig\_and\_dummy' does in fact have a dummy variable indicating whether a day is a weekend or not. The table showing coefficients and p-values are sorted from highest p-value to lowest p-value.

```

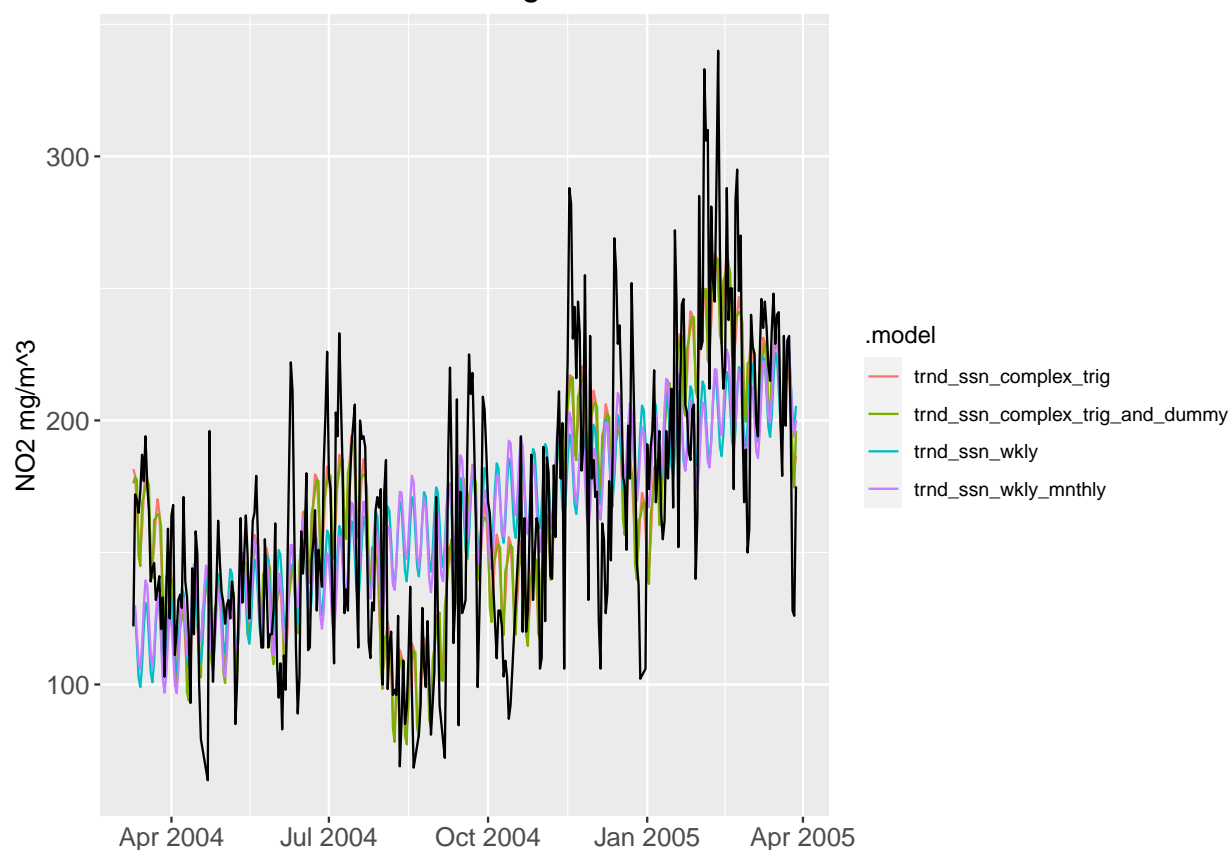
models %>% dplyr::select(trnd_ssn_wkly, trnd_ssn_wkly_mnthly,
                        trnd_ssn_complex_trig,
                        trnd_ssn_complex_trig_and_dummy) %>% coef() %>%
                        arrange(desc(p.value))

## # A tibble: 43 x 6
##   .model          term          estimate std.error statistic p.value
##   <chr>          <chr>          <dbl>    <dbl>    <dbl>    <dbl>
## 1 trnd_ssn_complex_tr~ sin(2 * pi * step/~ 0.00797      2.64    0.00302 0.998
## 2 trnd_ssn_complex_tr~ sin(2 * pi * step/~ -0.0453      2.59   -0.0175 0.986
## 3 trnd_ssn_wkly_mnthly cos(2 * pi * step/~ -0.931      3.05   -0.305 0.761
## 4 trnd_ssn_complex_tr~ cos(2 * pi * step/~ -2.09      3.11   -0.672 0.502
## 5 trnd_ssn_complex_tr~ sin(2 * pi * step/~ 4.44      3.46    1.28 0.200
## 6 trnd_ssn_complex_tr~ sin(2 * pi * step/~ 3.33      2.44    1.36 0.173
## 7 trnd_ssn_complex_tr~ sin(2 * pi * step/~ 3.33      2.39    1.39 0.165
## 8 trnd_ssn_complex_tr~ sin(2 * pi * step/~ 4.85      2.58    1.88 0.0604
## 9 trnd_ssn_wkly      cos(2 * pi * step/~ 5.84      3.09    1.89 0.0597
## 10 trnd_ssn_complex_tr~ sin(2 * pi * step/~ 4.81      2.53    1.90 0.0578
## # ... with 33 more rows

models %>%
  dplyr::select(trnd_ssn_wkly, trnd_ssn_wkly_mnthly,
                trnd_ssn_complex_trig, trnd_ssn_complex_trig_and_dummy) %>%
  fitted() %>% autoplot() +
  autolayer(dailyAQ_train) +
  ggtitle("Trend-Season Model with Trigonometric Functions") +
  ylab("NO2 mg/m^3") + xlab("") + my_theme
## Plot variable not specified, automatically selected '.vars = .fitted'
## Plot variable not specified, automatically selected '.vars = NO2'

```

## Trend–Season Model with Trigonometric Functions



Visually, the complex trig functions track the process very well though there are a few insignificant variables. Interestingly, the cosine function counterpart to the most insignificant variable of the `trnd_ssn_complex_trig` model (the sine function component of a 76.8 day season) has a p-value significant to the 0.001 level. This is not shown in the table but could easily be recreated from the attached `.rmd`.

To decide which is the best approach, we compared performance metrics across the candidate models with a trend and seasonal component.

```
models %>%
  dplyr::select(trnd_ssn_wkly, trnd_ssn_wkly_mnthly, trnd_ssn_complex_trig,
    trnd_ssn_complex_trig_and_dummy, trnd_ssn_w_kday_dummy, trnd_ssn_w_knd_dummy) %>%
  report() %>% dplyr::select(.model, adj_r_squared, AIC, BIC, df.residual)
## Warning in report.mdl_df(.): Model reporting is only supported for individual
## models, so a glance will be shown. To see the report for a specific model, use
## 'select()' and 'filter()' to identify a single model.
## # A tibble: 6 x 5
##   .model                                adj_r_squared  AIC    BIC df.residual
##   <chr>                                <dbl> <dbl> <dbl>    <int>
## 1 trnd_ssn_w_kday_dummy                0.352 2887. 2922.     376
## 2 trnd_ssn_w_knd_dummy                 0.353 2882. 2897.     381
## 3 trnd_ssn_wkly                       0.340 2890. 2910.     380
## 4 trnd_ssn_wkly_mnthly                 0.350 2886. 2914.     378
## 5 trnd_ssn_complex_trig                0.594 2715. 2782.     368
## 6 trnd_ssn_complex_trig_and_dummy      0.610 2701. 2772.     367
```


The trend/season models with the complex trigonometric functions clearly outperformed the other models, but we decided to include other model candidates in case these are overfit. Since the only difference between the `trnd_ssn_wkly` and `trnd_ssn_wkly_mnthly` model is the inclusion of a monthly trigonometric function and the latter performed better, we decided to exclude the `trnd_ssn_wkly` model from the rest of our analysis.

## 2.2 Modeling the residuals of best trend/season models

Next, we extracted and plotted the residuals of each candidate trend/season model to inspect for autoregressive properties.

```
# Extract model residuals
data_resids <- models %>%
  dplyr::select(trnd_ssn_wkly_mnthly, trnd_ssn_w_wkday_dummy,
               trnd_ssn_w_wknd_dummy, trnd_ssn_complex_trig,
               trnd_ssn_complex_trig_and_dummy) %>%
  residuals()


# Plot the residuals
data_resids %>% autoplot(.resid) + ggtitle("Trend/Season Model Residuals") +
  xlab("") + my_theme
```




There is clearly some autocorrelation in the residuals in all of the models because none of them look like random white noise. The next step was to inspect the ACF and PACF plots to decide how to model them. Because there are so many models, we look at the dummy variable models separate from the trigonometric models.

```
plt_resid_ACF <- data_resids %>%  
  dplyr::filter(.model %in% c("trnd_ssn_w_wknd_dummy" ,"trnd_ssn_w_wkday_dummy")) %>%  
  ACF(.resid) %>% autoplot() + ggtitle("Dummies: NO2 Residuals ACF plot") + my_theme  
  
plt_resid_PACF <- data_resids %>%  
  dplyr::filter(.model %in% c("trnd_ssn_w_wknd_dummy" ,"trnd_ssn_w_wkday_dummy")) %>%  
  PACF(.resid) %>% autoplot() + ggtitle("Dummies: NO2 Residuals PACF plot") + my_theme  
  
ggarrange(plt_resid_ACF,plt_resid_PACF,nrow=2,ncol=1)
```

Dummies: NO2 Residuals ACF plot



Dummies: NO2 Residuals PACF plot



Both dummy variable trend/season models exhibit similar behavior in the ACF and PACF. There is sinusoidal decay in the ACF of the residuals. The PACF could be interpreted as cutting off after lag 2 or as exhibiting sinusoidal behavior. This indicates that the residuals could be modeled as either AR(2) or as an ARMA process.


```
plt_resid_ACF <- data_resids %>%
  dplyr::filter(.model %in% c("trnd_ssn_wkly_mnthly" , "trnd_ssn_complex_trig",
    "trnd_ssn_complex_trig_and_dummy")) %>%
  ACF(.resid) %>%
  autoplot() +
  ggtitle("Trig Functions: NO2 Residuals ACF plot") +
  my_theme

plt_resid_PACF <- data_resids %>%
  dplyr::filter(.model %in% c("trnd_ssn_wkly_mnthly" , "trnd_ssn_complex_trig",
    "trnd_ssn_complex_trig_and_dummy")) %>%
  PACF(.resid) %>% autoplot() +
  ggtitle("Trig Functions: NO2 Residuals PACF plot") +


ggarrange(plt_resid_ACF, plt_resid_PACF, nrow=2, ncol=1)
```



Trig Functions: NO2 Residuals ACF plot



Trig Functions: NO2 Residuals PACF plot



The two trend/season models with more trigonometric functions (denoted ‘complex’) potentially show sinusoidal behavior in the ACF and the PACF and the PACF appears to cut off after 1 lag, although the second lag is technically barely above the significant threshold for the one that includes a weekend dummy variable. This indicates that it these should either be modeled as an AR(1) process or possibly an ARMA process if the PACF could be considered to be sinusoidal. The less complex trend/season model that only has monthly and weekly components looks more like the dummy variable models above with sinusoidal decay in the ACF and a PACF that either cuts off at lag 2, indicating AR(2), or a sinusoidal PACF which would lead to modeling it as an ARMA process.

Next we performed autoregressive modeling of the residuals according to these hypotheses. Because of the structure of the tibble package, we had to apply the AR(1) and AR(2) models to all residual datasets. We also deployed an auto-selection process.

```
# Model the residuals
models_residuals <- data_resids %>%
  rename("model" = ".model", "resid" = ".resid") %>%
  model(auto.arima = ARIMA(resid ~ PDQ(0,0,0)),
        ar2 = ARIMA(resid ~ pdq(2, 0, 0) + PDQ(0,0,0)),
        ar1 = ARIMA(resid ~ pdq(1, 0, 0) + PDQ(0,0,0)))

# Show the pdq values from the auto.arima models
models_residuals %>% tidy() %>% filter(.model == 'auto.arima') %>% print(n = 3e3)
## # A tibble: 11 x 7
##   model                .model term estimate std.error statistic p.value
##   <chr>                <chr>  <chr>   <dbl>    <dbl>    <dbl>   <dbl>
## 1 trnd_ssn_complex_trig auto.ar~ ar1    0.403    0.0468     8.61 1.85e-16
## 2 trnd_ssn_complex_trig_a auto.ar~ ar1    0.431    0.0462     9.33 8.34e-19
## 3 trnd_ssn_w_wkday_dummy auto.ar~ ar1    1.33     0.144     9.22 1.97e-18
## 4 trnd_ssn_w_wkday_dummy auto.ar~ ar2   -0.375    0.115    -3.25 1.25e- 3
## 5 trnd_ssn_w_wkday_dummy auto.ar~ ma1   -0.763    0.124    -6.17 1.74e- 9
## 6 trnd_ssn_w_wknd_dummy auto.ar~ ar1    1.32     0.139     9.49 2.52e-19
## 7 trnd_ssn_w_wknd_dummy auto.ar~ ar2   -0.369    0.111    -3.31 1.01e- 3
## 8 trnd_ssn_w_wknd_dummy auto.ar~ ma1   -0.764    0.119    -6.41 4.39e-10
## 9 trnd_ssn_wkly_mnthly auto.ar~ ar1    1.29     0.121    10.7 2.03e-23
## 10 trnd_ssn_wkly_mnthly auto.ar~ ar2   -0.337    0.0972    -3.47 5.78e- 4
## 11 trnd_ssn_wkly_mnthly auto.ar~ ma1   -0.772    0.101    -7.65 1.66e-13
```

The auto-selection process yielded ARMA(2, 1) models for the day-of-the-week and weekend/not weekend dummy variable models as well as the weekly and monthly trigonometric models. The auto-selection process yielded an AR(1) model for the complex\_trig models as we hypothesized based on the ACF and PACF.

Next we report the AIC and BIC for each of the candidate residual models.

```
# Compare AIC and BIC across the models; sort by AIC
models_residuals %>% report() %>%
  dplyr::select(model, .model, AIC, BIC) %>% arrange(AIC) %>% print(n = 100)
## Warning in report.mdl_df(.): Model reporting is only supported for individual
## models, so a glance will be shown. To see the report for a specific model, use
## 'select()' and 'filter()' to identify a single model.
## # A tibble: 15 x 4
##   model                .model      AIC    BIC
##   <chr>                <chr>    <dbl> <dbl>
## 1 trnd_ssn_complex_trig_and_dummy auto.arima 3680. 3688.
## 2 trnd_ssn_complex_trig_and_dummy ar1        3680. 3688.
## 3 trnd_ssn_complex_trig_and_dummy ar2        3682. 3694.
```

|       |                        |            |       |       |
|-------|------------------------|------------|-------|-------|
| ## 4  | trnd_ssn_complex_trig  | auto.arima | 3707. | 3715. |
| ## 5  | trnd_ssn_complex_trig  | ar1        | 3707. | 3715. |
| ## 6  | trnd_ssn_complex_trig  | ar2        | 3709. | 3721. |
| ## 7  | trnd_ssn_w_wkday_dummy | auto.arima | 3723. | 3739. |
| ## 8  | trnd_ssn_w_wkday_dummy | ar2        | 3725. | 3737. |
| ## 9  | trnd_ssn_w_wkday_dummy | ar1        | 3732. | 3740. |
| ## 10 | trnd_ssn_w_wknd_dummy  | auto.arima | 3734. | 3750. |
| ## 11 | trnd_ssn_w_wknd_dummy  | ar2        | 3737. | 3748. |
| ## 12 | trnd_ssn_w_wknd_dummy  | ar1        | 3743. | 3751. |
| ## 13 | trnd_ssn_wkly_mnthly   | auto.arima | 3754. | 3770. |
| ## 14 | trnd_ssn_wkly_mnthly   | ar2        | 3761. | 3773. |
| ## 15 | trnd_ssn_wkly_mnthly   | ar1        | 3766. | 3774. |

Based on AIC, the auto-selected ARIMA model performed best or tied for the best model for each set of residuals. BIC did not align exactly with AIC results. For both the dummy variable trend/season model residuals, the AR(2) model had the best score which was our hypothesis based on the AIC and BIC. The score improvement is likely negligible. For ease of visualization, we decided to continue with the rest of our analysis using only the auto-selected models.

Next we plotted residuals versus fitted and QQ plots of model residuals to test for the iid assumption.


```
# Choose the auto.arima model for the rest of the analysis
models_residuals <- models_residuals %>% dplyr::select(-ar2, -ar1)

# Plot residuals vs. fitted
models_residuals_fitted <-
  models_residuals %>% fitted() %>% as_tibble() %>% dplyr::select(Date, '.fitted')

models_residuals_resid <-
  models_residuals %>% residuals() %>% as_tibble()


model_residuals_fittedAndResiduals <-
  inner_join(models_residuals_resid, models_residuals_fitted, by = "Date")

ggplot(model_residuals_fittedAndResiduals, aes(.fitted, .resid)) +
  geom_point() + facet_grid(rows = vars(model), cols = vars(.model))
```



There is no clear trend in any of the model residuals.

```
# Plot QQ plots  
ggplot(model_residuals_fittedAndResiduals, aes(sample = .resid)) +  
  stat_qq() + stat_qq_line(color = "red") +  
  facet_grid(rows = vars(model), cols = vars(.model))
```



The residuals for each model all appear to exhibit similar patterns and are approximately gaussian with some fat-tail behavior.

The next step was to forecast and test each model against the test set.

### 3 Forecasting

We started by forecasting the residuals.


```
resid_summary <- models_residuals %>% augment()

resid_7day_forecast <- models_residuals %>%
  forecast(h = "1 week")

plt_resid_forecast <- resid_7day_forecast %>% feasts::autoplot(alpha = 0.8) +
  geom_line(data = resid_summary, aes(x = Date, y = resid)) +
  ggtitle("Forecast of residuals") + my_theme

plt_resid_forecast
```

## Forecast of residuals





```
resid_7day_forecast <- resid_7day_forecast %>% hilo(level = 95)
```

The forecast of residuals look reasonable and very similar as does the width of the confidence intervals. Next we forecasted the trend/season models.


```
main_summary <- models %>% select(trnd_ssn_w_wknd_dummy,
                                trnd_ssn_w_wkday_dummy,
                                trnd_ssn_complex_trig_and_dummy,
                                trnd_ssn_complex_trig,
                                trnd_ssn_wkly_mnthly) %>% augment()

main_7day_forecast <- models %>%
  select(trnd_ssn_w_wknd_dummy, trnd_ssn_w_wkday_dummy,
         trnd_ssn_complex_trig_and_dummy,
         trnd_ssn_complex_trig, trnd_ssn_wkly_mnthly) %>%
  forecast(new_data = dailyAQ_test)

plt_main_forecast <- main_7day_forecast %>% feasts::autoplot(alpha = 0.8) +
  geom_line(data = main_summary, aes(x = Date, y = NO2)) +
  ggtitle("Forecast of Trend-Season Model") +
  facet_grid(rows = vars(.model)) +
  my_theme

plt_main_forecast
```

# Forecast of Trend–Season Model



The trend/season forecasts also look similar with similar confidence intervals. The next step was to combine these models to form the full forecast and to evaluate performance. We show the mean squared error for each model after the plot below.

### 3.1 Combining Season/trend and residuals forecasts and computing MSE

```
full_model_summary <- resid_summary %>%
  left_join(main_summary,
    by = c("Date" = "Date", "model" = ".model"),
    suffix = c(".resid", ".full"))

full_forecast <- left_join(resid_7day_forecast, main_7day_forecast,
  by = c("model" = ".model", "Date" = "Date"),
  suffix = c(".resid", ".full")) %>%
  mutate(pt_pred = .mean.resid + .mean.full)

test_data <- dailyAQ_test %>%
  dplyr::select(Date, "NO2") %>%
  rename("NO2_Actual" = "NO2")

full_forecast <- full_forecast %>% inner_join(test_data, by = ("Date" = "Date"))

forecast_MSE <- full_forecast %>% as_tibble() %>% group_by(model, .model) %>%
  mutate(sqrd_error = (pt_pred - NO2_Actual)^2) %>%
  summarize(MSE = mean(sqrd_error))
## 'summarise()' regrouping output by 'model' (override with '.groups' argument)
```


### 3.2 Plotting forecasts and MSE

```
full_forecast <- full_forecast %>% unpack_hilo(cols = "95%") %>%
  mutate('95%_upper' = .mean.full + '95%_upper') %>%
  mutate('95%_lower' = .mean.full + '95%_lower') %>%
  mutate(model_ID = paste(model, .model, sep = "_"))

plot_forecasts <- ggplot(data = full_forecast, aes(group = model_ID, color = model_ID)) +
  geom_line(aes(x=Date,y=NO2_Actual), color = "black") +
  geom_line(aes(x=Date,y=pt_pred)) +
  geom_line(aes(x=Date,y='95%_lower'), linetype="dashed") +
  geom_line(aes(x=Date,y='95%_upper'), linetype="dashed") +
  xlab("") + ylab("NO2 mg/m^3") +
  ggtitle("NO2 Season/Trend Model + ARMA of Residuals with 95% CIs") + my_theme

plot_forecasts
```

## NO2 Season/Trend Model + ARMA of Residuals with 95% CIs



```
forecast_MSE
## # A tibble: 5 x 3
## # Groups:   model [5]
##   model                                .model      MSE
##   <chr>                                <chr>      <dbl>
## 1 trnd_ssn_complex_trig                auto.arima  346.
## 2 trnd_ssn_complex_trig_and_dummy      auto.arima  405.
## 3 trnd_ssn_w_wkday_dummy               auto.arima  562.
## 4 trnd_ssn_w_wknd_dummy               auto.arima  448.
## 5 trnd_ssn_wkly_mnthly                 auto.arima  324.
```

Despite being significantly outperformed based on AIC, BIC, adjust  $R^2$ , and residuals, the model with the trigonometric functions representing weekly and monthly cycles performed best on the test dataset. That is the model we have selected. This indicates that the models with many trigonometric functions were overfitted.

Next we used the best model to simulate the next year of data.

## 4 Simulation

### 4.1 Simulating one year proceeding our observations

We started by forecasting the residuals.

```
resid_1year_simulation <- models_residuals %>% dplyr::select(auto.arima) %>%
  filter(model == "trnd_ssn_wkly_mnthly") %>%
  forecast(h = "1 year") # %>% hilo(level = 95)

resid_1year_simulation <- models_residuals %>% dplyr::select(auto.arima) %>%
  filter(model == "trnd_ssn_wkly_mnthly") %>% generate(h = "1 year", seed = 1)
```

Next we forecasted the best trend/season model.

```
vec_simulation_dates <- seq(ymd('2005-03-29'), ymd('2006-03-29'), by='days')

original_steps <- length(dailyAQ_train$Date)
sim_steps <- length(vec_simulation_dates)

simulation_dates <- as_tibble() %>% as_tibble(.rows = length(vec_simulation_dates)) %>%
  bind_cols(vec_simulation_dates) %>% rename('Date' = '...1') %>%
  mutate(NO2 = -999) %>%
  mutate(step = (original_steps+1):(original_steps + sim_steps)) %>%
  mutate(day = as.factor(as.character(lubridate::wday(Date, label = TRUE)))) %>%
  mutate(weekend = isWeekend(Date)) %>% as_tsibble(index = Date)
## Warning: The 'x' argument of 'as_tibble()' can't be missing as of lifecycle 3.0.0.
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_warnings()' to see where this warning was generated.
## New names:
## * NA -> ...1

main_1year_forecast <- models %>%
  select(trnd_ssn_wkly_mnthly) %>%
  forecast(new_data = simulation_dates)
```


Finally, we combined them into a single simulation.

```
full_simulation <- left_join(resid_1year_simulation, main_1year_forecast,
  by = c("model" = ".model", "Date" = "Date"),
  suffix = c(".resid", ".full")) %>%
  mutate(pt_pred = .sim + .mean) %>%
  mutate(model_ID = paste(model, ".model", "_"))

dailyAQ_train <- dailyAQ_train %>% mutate(model_ID = "Training Data")

plot_simulation <- ggplot(data = full_simulation, aes(group = model_ID,
  color = model_ID)) + geom_line(aes(x=Date, y=pt_pred)) +
  geom_line(data = dailyAQ_train, aes(x = Date, y = NO2)) +
  xlab("") + ylab("NO2 mg/m^3") +
  ggtitle("NO2 Season/Trend Model + ARMA of Residuals with 95% CIs") + my_theme

plot_simulation
```



The final simulation looks reasonable, so we are comfortable reporting this model as our final product. Next we compared the coefficient on the trend terms from the trend/season model built from the observations and a trend/season model built from the simulation.

## 4.2 Comparing trend/season models built from observations vs. simulation

```
# Isolate the final trend/season model used for simulation
models %>% dplyr::select(trnd_ssn_wkly_mnthly)
## # A tibble: 1 x 1
##   trnd_ssn_wkly_mnthly
##               <model>
## 1               <TSLM>

# Create a trend/season model of the simulation
simulation_model <- full_simulation %>% dplyr::select(-model, -.model, -.rep) %>%
  model(sim_trnd_ssn_wkly_mnthly = TSLM(pt_pred ~ trend() + sin(2*pi*step/(365.25/12)) +
    cos(2*pi*step/(365.25/12)) + sin(2*pi*step/7) + cos(2*pi*step/7)))

# Show the model coefficients and p-values
models %>% dplyr::select(trnd_ssn_wkly_mnthly) %>% coef()
## # A tibble: 6 x 6
##   .model          term      estimate std.error statistic  p.value
##   <chr>          <chr>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 trnd_ssn_wkly_mn~ (Intercept)    113.     4.34     26.1 1.33e-86
## 2 trnd_ssn_wkly_mn~ trend()         0.261    0.0195    13.4 1.42e-33
## 3 trnd_ssn_wkly_mn~ sin(2 * pi * step/(36~    8.77     3.07     2.86 4.54e- 3
## 4 trnd_ssn_wkly_mn~ cos(2 * pi * step/(36~   -0.931    3.05    -0.305 7.61e- 1
## 5 trnd_ssn_wkly_mn~ sin(2 * pi * step/7)    14.9     3.06     4.87 1.67e- 6
## 6 trnd_ssn_wkly_mn~ cos(2 * pi * step/7)     5.84     3.06     1.90 5.76e- 2
```

```
simulation_model %>% coef()
## # A tibble: 6 x 6
##   .model          term      estimate std.error statistic    p.value
##   <chr>          <chr>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 sim_trnd_ssn_wkly~ (Intercept)    219.      4.08      53.8 8.69e-174
## 2 sim_trnd_ssn_wkly~ trend()      0.262    0.0193     13.6 3.60e- 34
## 3 sim_trnd_ssn_wkly~ sin(2 * pi * step/(~    8.74      2.87      3.04 2.53e- 3
## 4 sim_trnd_ssn_wkly~ cos(2 * pi * step/(~   -6.16      2.87     -2.14 3.28e- 2
## 5 sim_trnd_ssn_wkly~ sin(2 * pi * step/7)    12.3      2.88      4.27 2.46e- 5
## 6 sim_trnd_ssn_wkly~ cos(2 * pi * step/7)     4.37      2.87      1.53 1.28e- 1

# Extract the trend coefficient
main_coefs_trend <- models %>% dplyr::select(trnd_ssn_wkly_mnthly) %>% coef() %>%
  filter(term == "trend()") %>% pull(estimate)
sim_coefs_trend <- simulation_model %>% coef() %>% filter(term == "trend()") %>%
  pull(estimate)

# Calculate and report the percent difference between the trend coefficients of
# the model built from observed data and the model built from the simulation
perc_dif <- (sim_coefs_trend - main_coefs_trend)/main_coefs_trend * 100


perc_dif
## [1] 0.5312944
```

The trend coefficient of the model built from the simulation is 0.53% higher than the trend coefficient of the model built from the observations. This is evidence that the simulation is sufficiently matching observed patterns.

Next we inspected the periodogram to verify that the simulation reproduces the seasonal patterns of the observed data.

### 4.3 Verifying simulation reproduces seasonality


```
N02_sim.ts <- ts(full_simulation %>% as_tibble() %>% dplyr::select(pt_pred))
pg.N02 <- spec.pgram(N02.ts, spans=9, demean=T, log='no')
```



```
pg.NO2_sim <- spec.pgram(NO2_sim.ts, spans=9, demean=T, log='no')
```



### Series: NO2\_sim.ts Smoothed Periodogram



The periodograms built from the observed data and the simulation data look very similar which further verifies the simulation. Next we compared sample statistics for the observed data and the simulated data.

```
obs_mean <- mean(dailyAQ_train$NO2)
obs_var <- var(dailyAQ_train$NO2)

sim_mean <- mean(full_simulation$pt_pred)
sim_var <- var(full_simulation$pt_pred)

obs_mean_percDiff <- (sim_mean - obs_mean)/obs_mean*100
obs_var_percDiff <- (sim_var - obs_var)/obs_var*100

obs_mean_percDiff
## [1] 63.2018
obs_var_percDiff
## [1] -12.94752
```

The percent difference between the observed mean and simulated mean is 63.2% higher which is unsurprising since the simulation assumes the positive trend observed in the observations continues into the simulation period. The variance is 12.9% lower in the simulation which is indicative of a possible limitation of the simulation: it might not capture extremes well.

Finally, to verify that the simulation captures the autocorrelation structure of the observations, we compared the ACF and PACF plots.

## 4.4 Verifying simulation reproduces autocorrelation structure

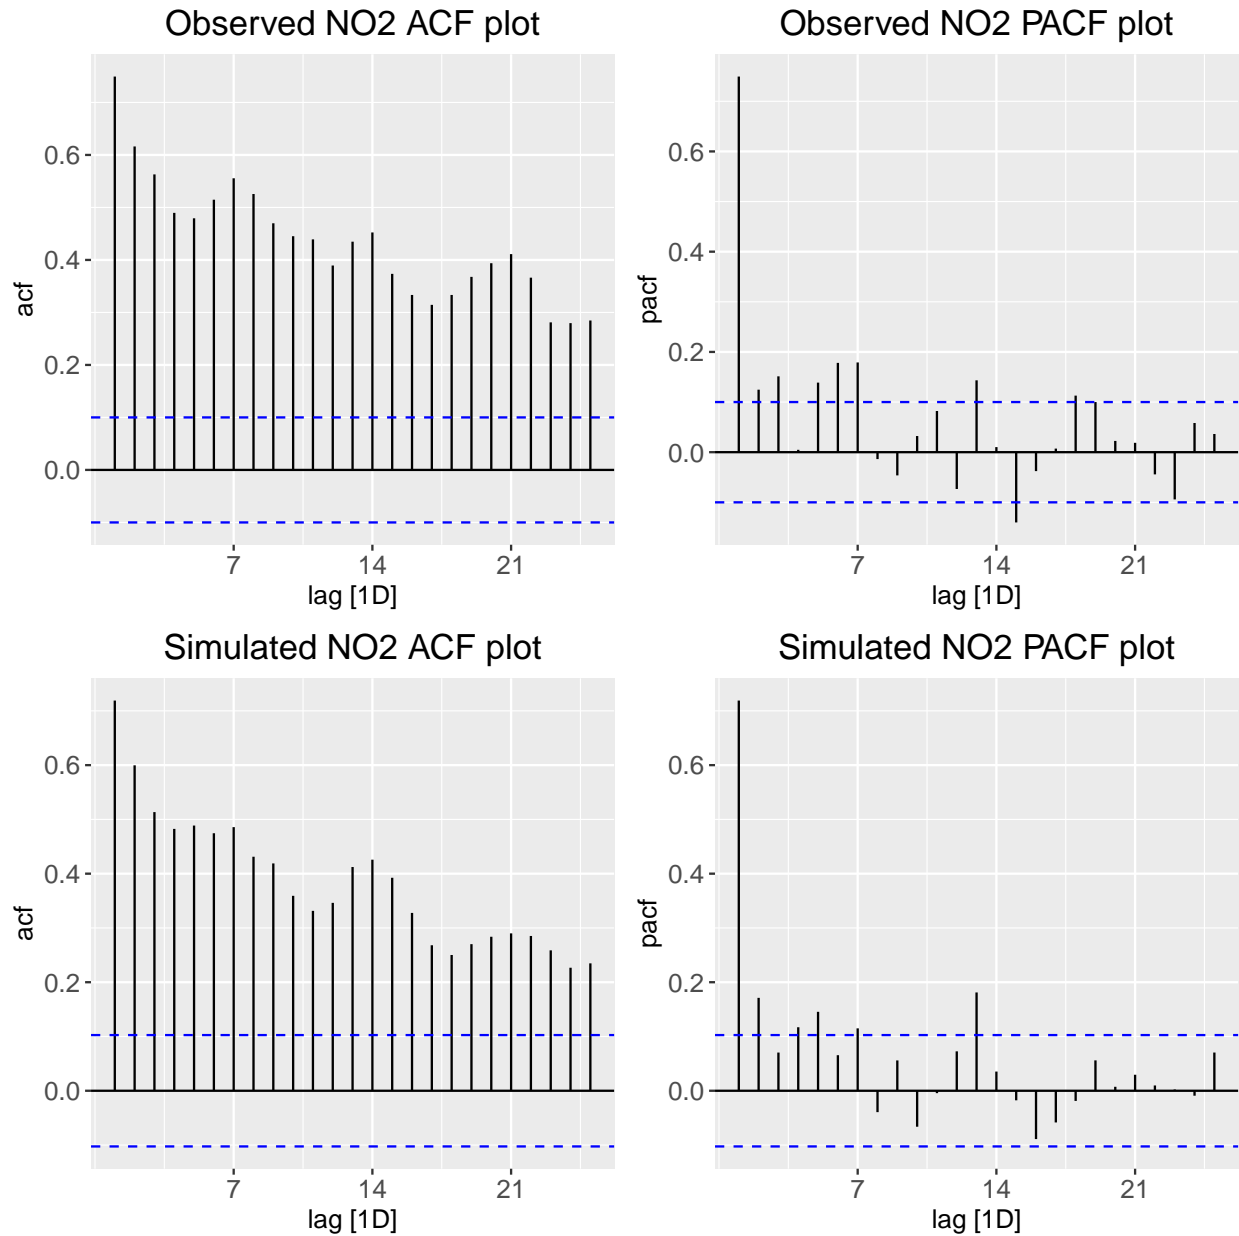
```
# Create plots of observed data
plt_N02_obs_ACF <- dailyAQ_train %>%
  ACF(N02) %>% autoplot() + ggtitle("Observed N02 ACF plot") + my_theme

plt_N02_obs_PACF <- dailyAQ_train %>%
  PACF(N02) %>% autoplot() + ggtitle("Observed N02 PACF plot") + my_theme

# Create plots of simulated data
plt_N02_sim_ACF <- full_simulation %>%
  ACF(pt_pred) %>% autoplot() + ggtitle("Simulated N02 ACF plot") + my_theme

plt_N02_sim_PACF <- full_simulation %>%
  PACF(pt_pred) %>% autoplot() + ggtitle("Simulated N02 PACF plot") + my_theme

ggarrange(plt_N02_obs_ACF, plt_N02_obs_PACF, plt_N02_sim_ACF, plt_N02_sim_PACF, nrow=2,
  ncol=2)
```



The ACF plot of both the observed data and the simulated data exhibit sinusoidal decay with similar levels of significance. The PACF plots show a steep drop off after the 1st lag with a few significant spikes and possible sinusoidal behavior. The similarities of these plots also serves to verify our simulation model.